

# Akash Verma

✉ as8052528735@gmail.com 📞 7905874934 📍 Noida, Uttar Pradesh 🌐 github.com/sky2d

## SUMMARY

Software developer with 2+ years of experience building scalable systems using **Node.js**, **Next.js**, and **FastAPI** — from distributed job pipelines to multi-provider AI integrations. Skilled in **backend engineering**, **system design**, and shipping production-ready applications with a focus on reliability and performance.

## EXPERIENCE

**ZnapAi** 🌐, *Software Developer*  
Full Time

02/2026 – 08/2026 | Delhi

- Maintained and scaled backend infrastructure, resolving critical **memory leaks** caused by base64-encoded API responses to support growing product load.
- Architected a multi-agent AI blogging pipeline capable of generating and auto-publishing **20+** SEO-optimized articles per day, designed to save users up to **80%** of content creation time
- Built an AI Playground with Next.js featuring **real-time streaming** responses, persistent session history, optimized media rendering, and cURL command extraction.

**AT Craya** 🌐, *Software Developer*  
Full Time

07/2024 – 01/2026 | Delhi

- Built a **video streaming pipeline** with encoding in **3 formats**, enabling adaptive playback and efficient media delivery at scale.
- Accelerated storefront launch timelines by **95%** (from an average of **10 days to 4 hours**) by engineering a self-serve drag-and-drop UI builder
- Built concurrent background sync engines with node-cron and the Instagram API, processing **2,500+ jobs/minute** without hitting API rate limits.

## PROJECTS

**Oshoot(AI Content Generation Platform)** 🌐, *Next.js 15, Node.js, PostgreSQL, Drizzle ORM, Inngest, OpenAI*

- Architected a declarative AI pipeline engine with fan-out execution, orchestrating **350+ concurrent jobs**; implemented exponential backoff retry logic that improved job success rate from **92% to 99.6%** during provider outages.
- Unified 5 AI providers (OpenAI, Anthropic, Gemini, fal.ai, ElevenLabs) into a single orchestration layer, scaling to **15,000+ multimedia assets/month** with average generation turnaround under 12 seconds per asset.
- Implemented S3 presigned URLs for direct client uploads, reducing Next.js server load by 65% optimized Drizzle ORM queries to maintain p95 latency under 180ms across **250,000+ user job records**.
- Deployed the platform on AWS ECS with S3 for media storage and Redis for caching session and job state.

**JobFlow** 🌐, *Next.js 15, Node.js, PostgreSQL, Kafka, Prisma*

- Architected an **event-driven background job processing system** using Node.js, Express.js, and Apache Kafka to decouple API requests from asynchronous workloads.
- Load tested with **Artillery** at **50 req/sec**, successfully ingesting **3,600+ jobs** in 2 minutes with a median **API latency of 10ms** analyzed saturation behavior via p95 latency tracking under peak load.
- Engineered fault tolerance via Kafka consumer retry policies and **Dead Letter Queue (DLQ)** mechanisms, achieving near **100% message** delivery success during sustained stress conditions.

## TECHNICAL SKILLS

**Frameworks** — React, React Native, Next.js, Node.js, Express.js, FastAPI, Tailwind CSS

**Languages** — JavaScript, TypeScript, Python, C/C++, HTML/CSS

**Databases** — PostgreSQL, Prisma, Drizzle, SQL

**Tools & Services** — Kafka, Git, Docker, Postman, Inngest

## EDUCATION

**Ajay Kumar Garg Engineering College**  
*Bachelor of Technology in Information Technology*

11/2021 – 11/2025 | Ghaziabad